

# Machine Learning for Marine Scientists

---

## Part 2: Data Preprocessing

29.09.2021

**Frauke Albrecht, Caroline Arnold, Jakob Lüttgau, Tobias Weigel**

Helmholtz AI Consultants for Earth and Environment

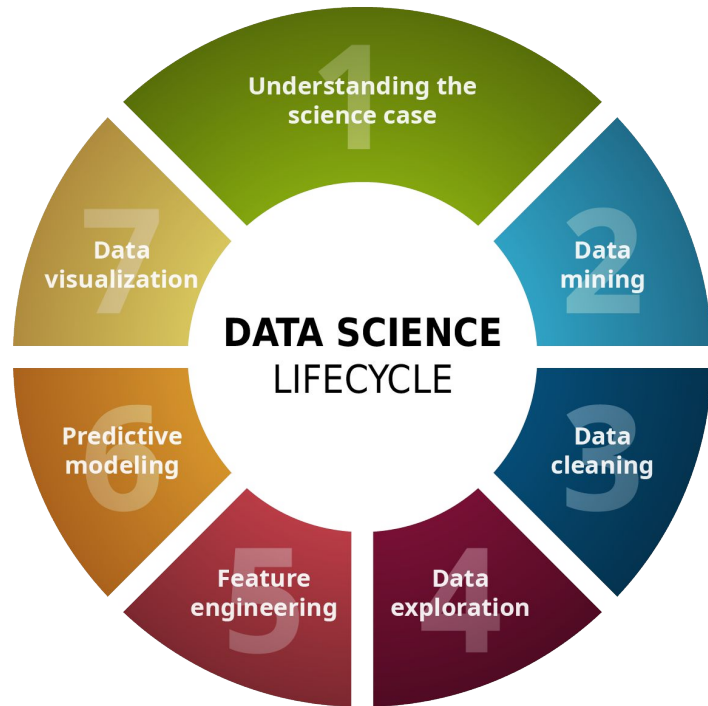
# Logistics

---

- Virtual room
  - Please share your cameras :-)
  - Please post questions in the shared google doc
  - Breakouts for the tutorial part in the afternoon
- Timeline for today
  - 10:00-12:00 - Lecture
  - 13:00-15:15 - Tutorial with jupyterlab

# Motivation

- Data is the most important resource in machine learning
- Scientific cases often use specialised data sets → specialised procedures
- Specialized hardware
- Scope today:
  - Everything that happens between raw data collection and the first ML training loop
  - Performance monitoring and improvement



# From raw data to training data

---

## Remote Sensing of Ocean Wind Speed

- 90% of observations cannot be used
- 60% of available features are not used

→ Filter the good samples and keep the interesting features

## Satellite Image Super-Resolution

- Find and remove cloudy scenes
- Annotate samples (expensive calculation)

→ Annotate the filtered samples

## Atmospheric Chemistry

- Very large number of small samples
- Cannot be shuffled in-memory

→ Shuffle the samples beforehand and save them in an accessible file format

## Preprocessing

- Clearly separated from ML training
- Executed once - before ML optimization
- Save new train / validation / test

# Modules in ML software

- Organize your code in a standardized way

```
cygnss_202003
|--- preprocessing
|----- preprocessing.py
|----- analysis.py
|--- training
|----- Model.py
|--- utils
|----- mathematics.py
```

- ❖ Helps with documentation and readability
- ❖ Helps with debugging
- ❖ Separate tasks in separate programs - if you add new data, know which parts of the code you need to touch
- ❖ We recommend a **preprocessing** module in any ML project

```
from cygnss_202003.preprocessing import preprocessing as prp

prp.open_dataset(...)
```

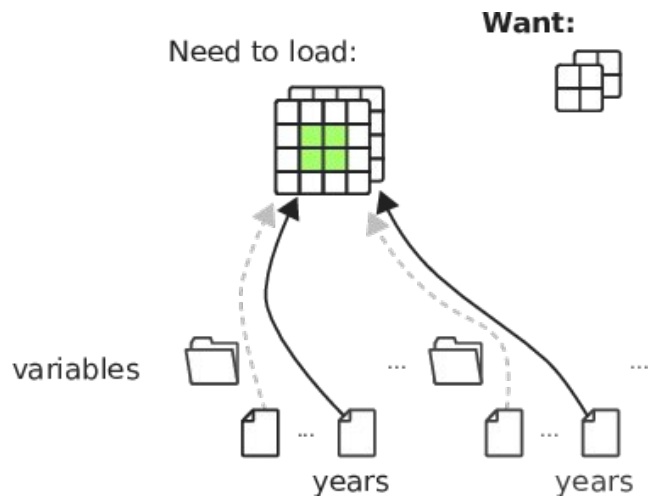
---

# A First Overview

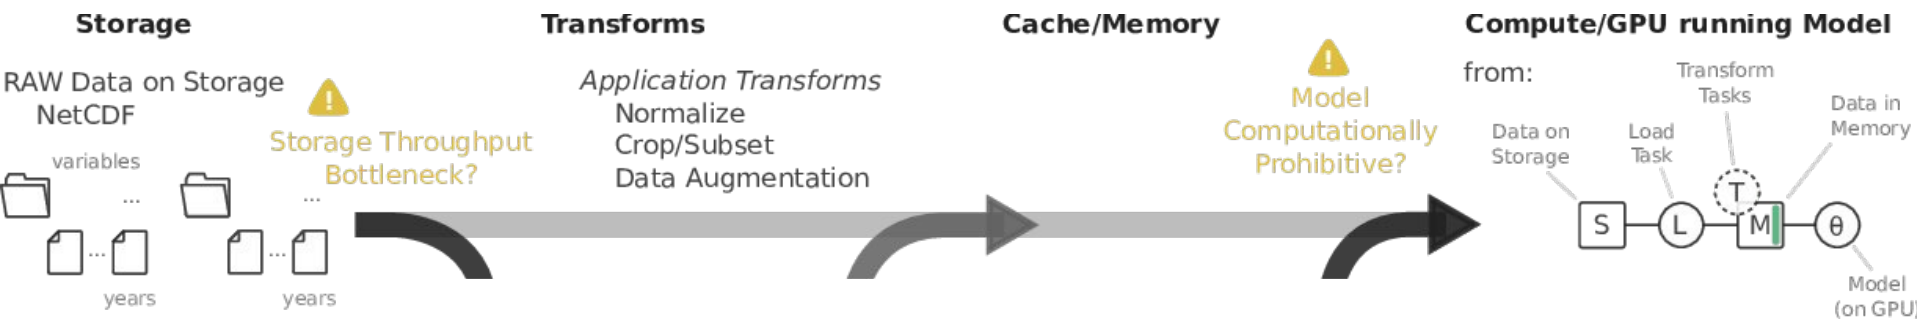
A workflow perspective

# The Basic Problem for many Pre-Processing Tasks

---

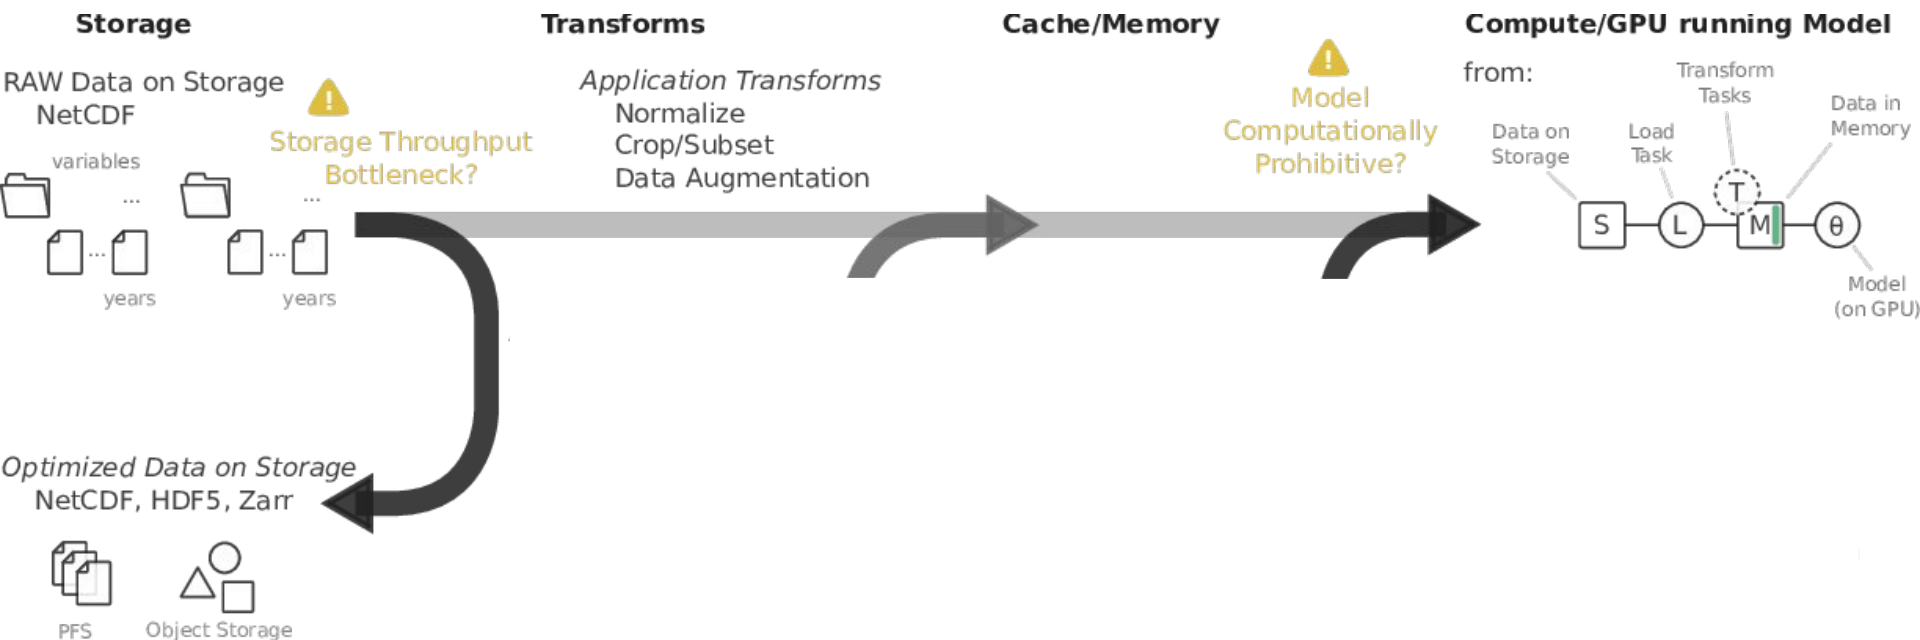


# A look at your end-to-end ML Pre-Processing Workflow:

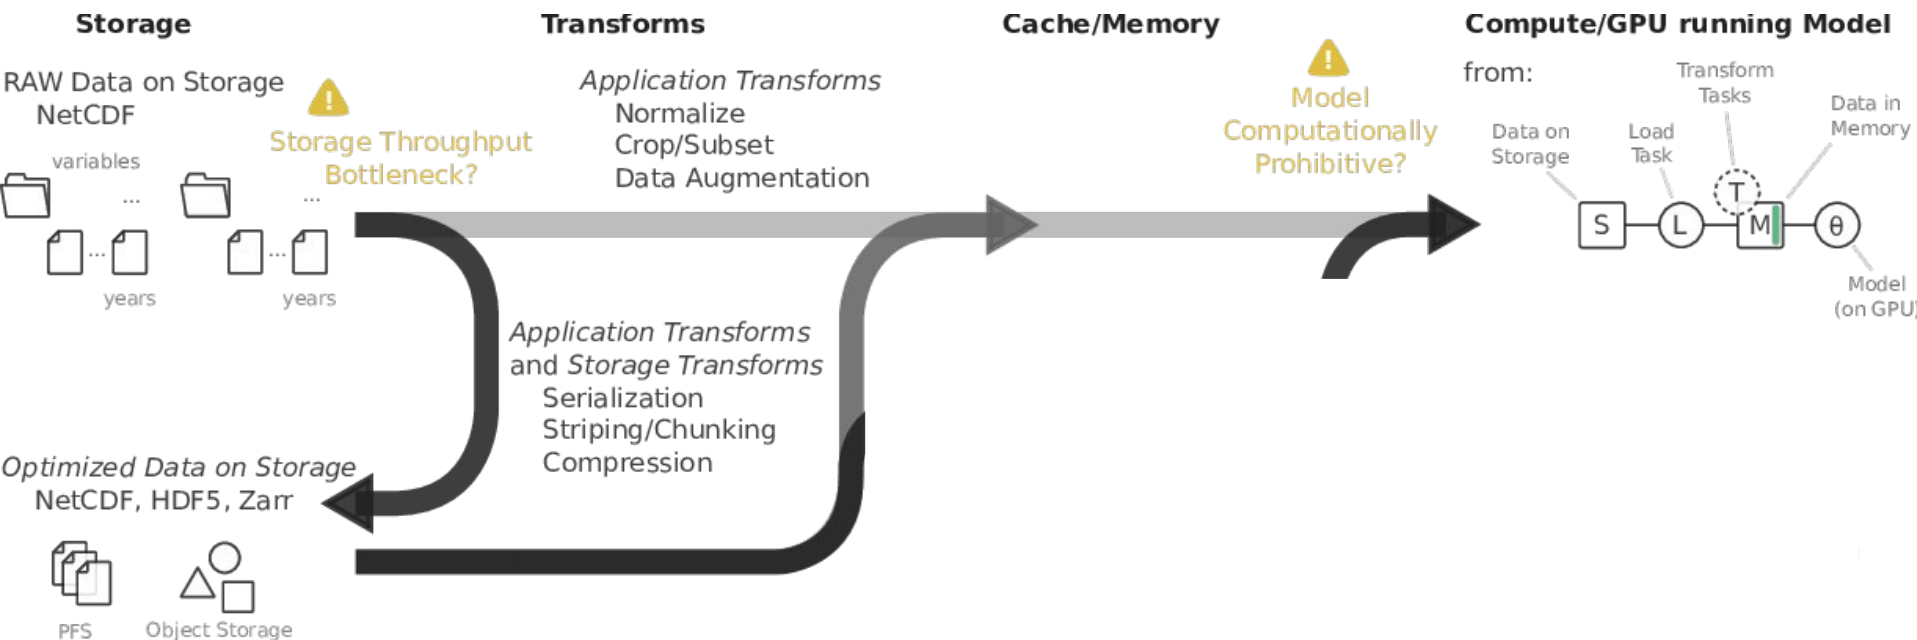




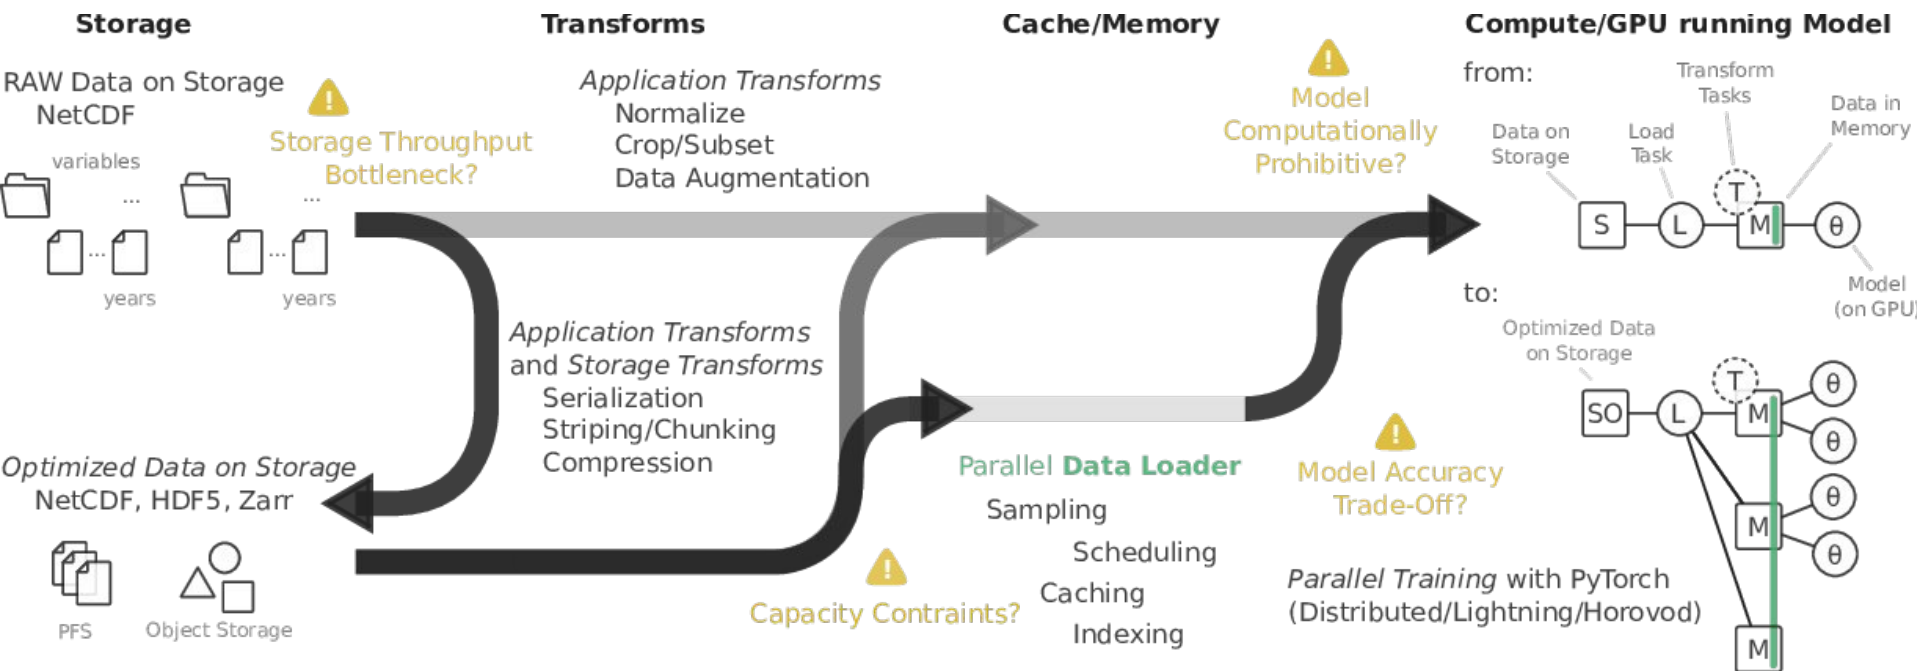
# A look at your end-to-end ML Pre-Processing Workflow:



# A look at your end-to-end ML Pre-Processing Workflow:



# A look at your end-to-end ML Pre-Processing Workflow:



---

# **HPC Systems: A look behind the curtain**

Memory Hierarchy, Cluster Computers, Storage Systems

# What happens when you load your data?

---

```
with open("/path/to/data", "r") as f:  
    data = f.read()
```

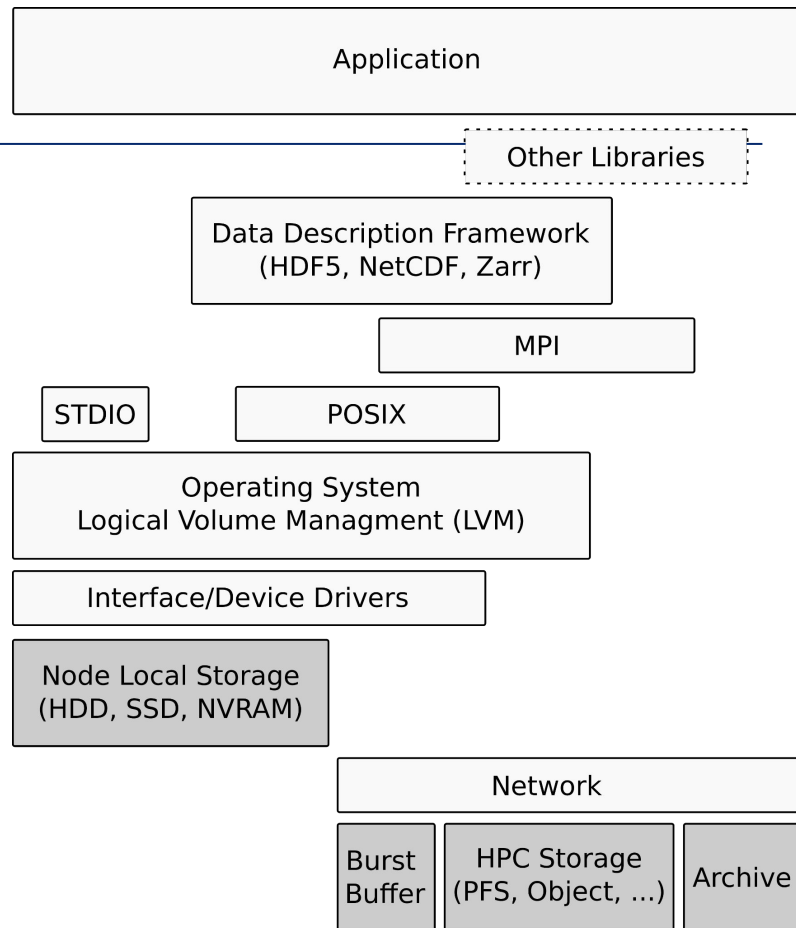
```
import pandas as pd  
df = pd.read_csv("/path/to/data.csv")
```

```
import h5py  
f = h5py.File('dataset.hdf5', 'r')
```

```
from netCDF4 import Dataset  
rootgrp = Dataset("/path/to/dataset.nc", "r", format="NETCDF4")
```

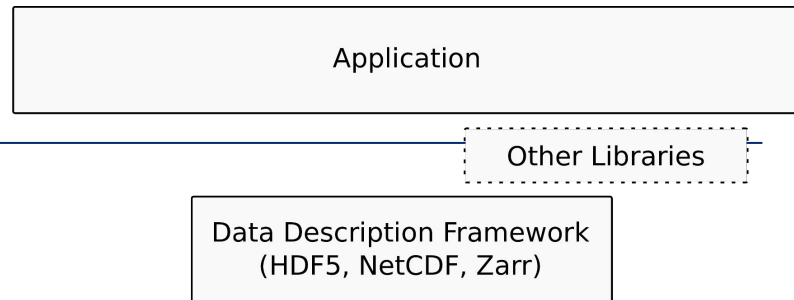
```
import xarray as xr  
ds = xr.open_dataset("/path/to/dataset.nc")
```

# System: Software Stack



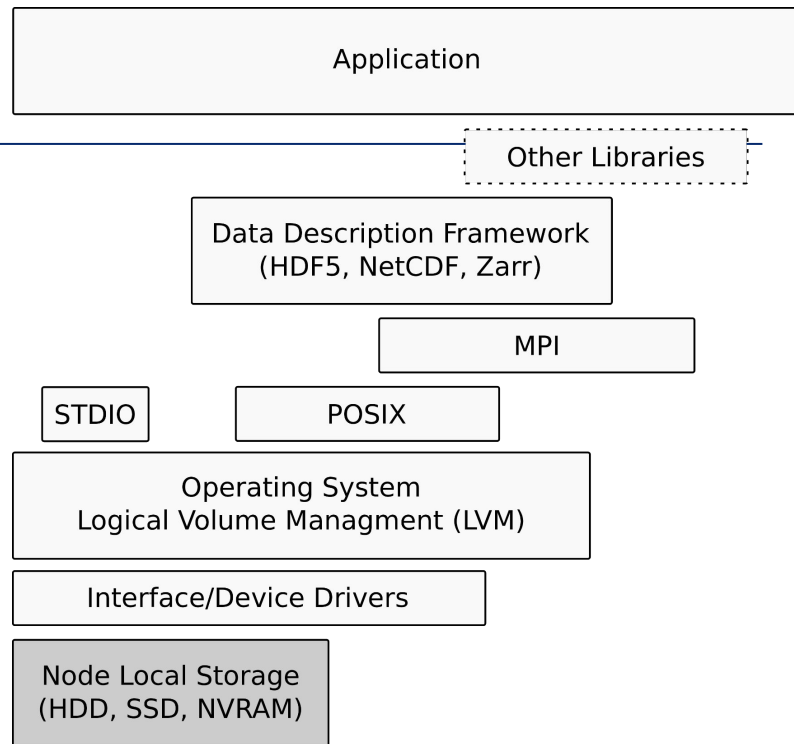
# System: Software Stack

---



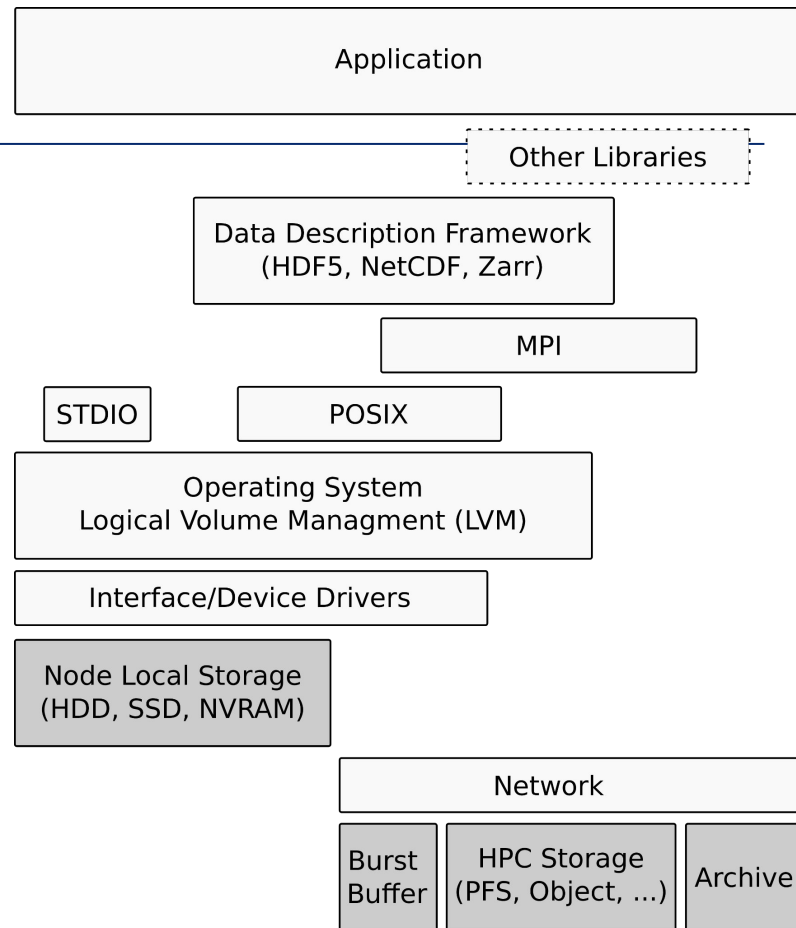
# System: Software Stack

---

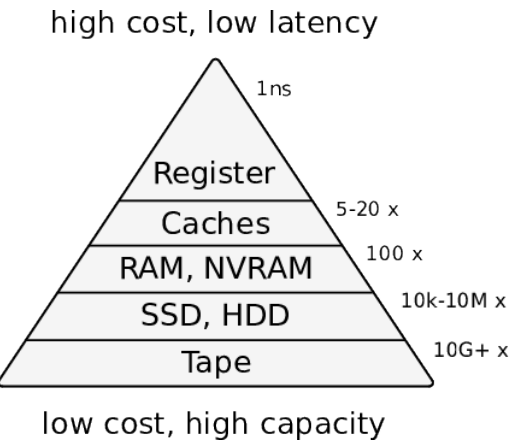




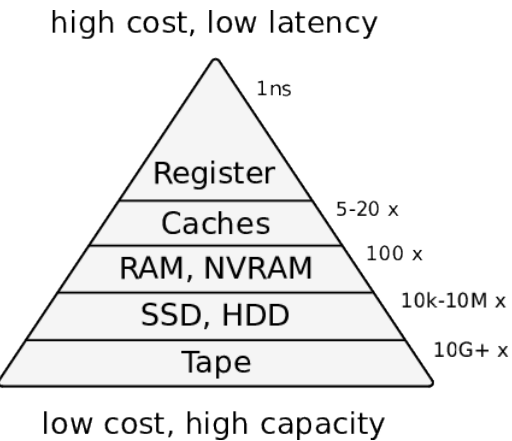
# System: Software Stack



# Memory Hierarchy



# Memory Hierarchy

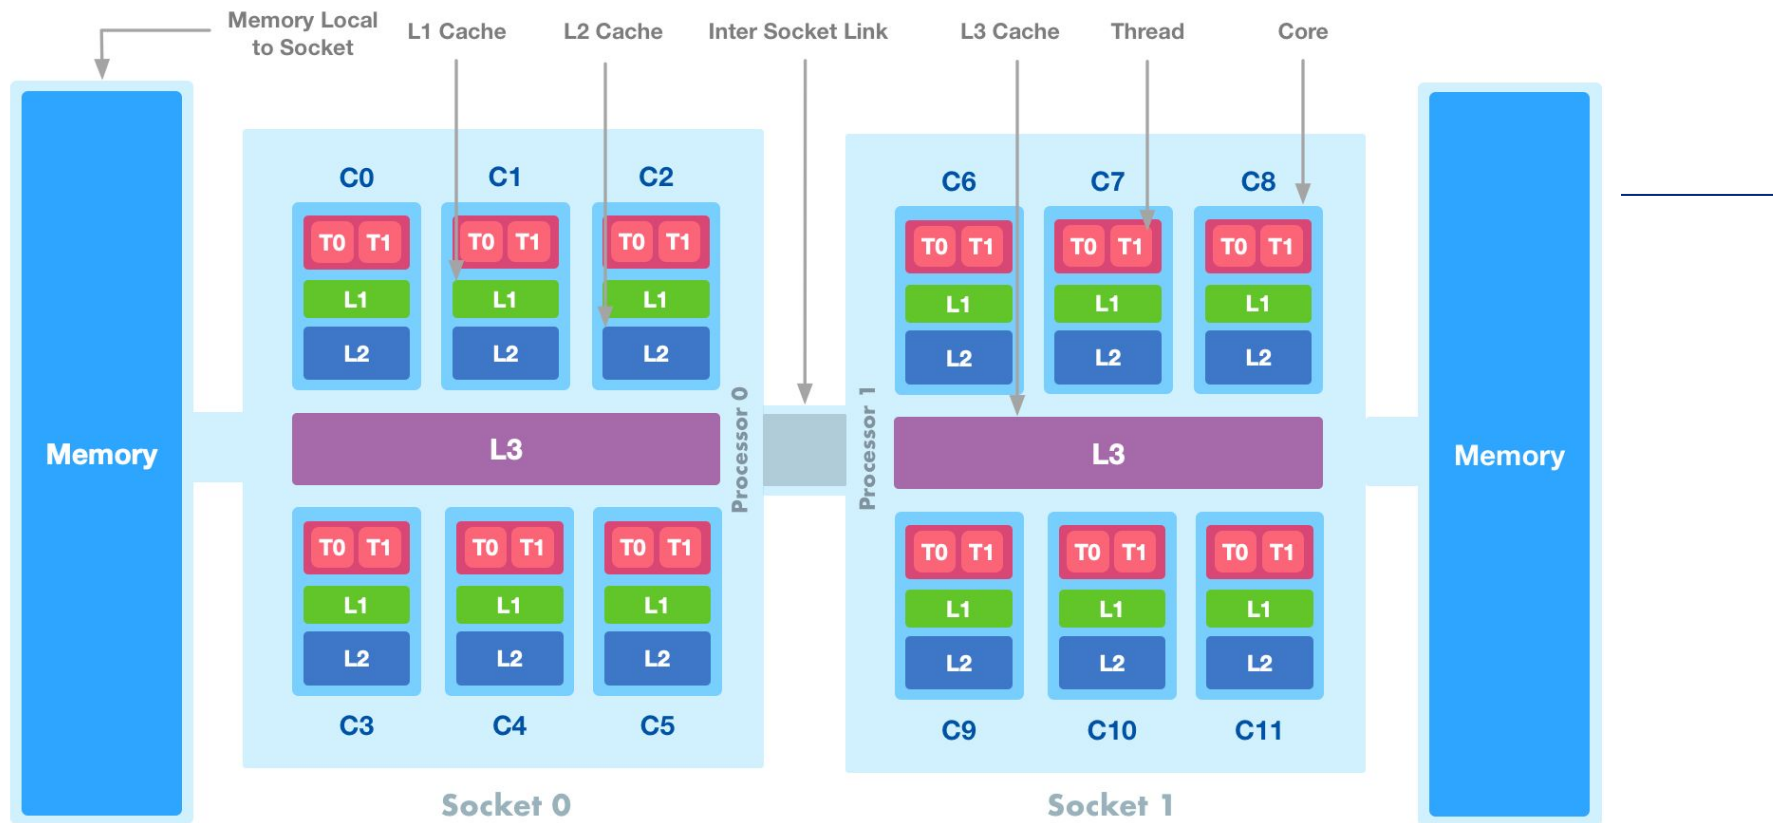


## The best case: Your data fits in memory

---

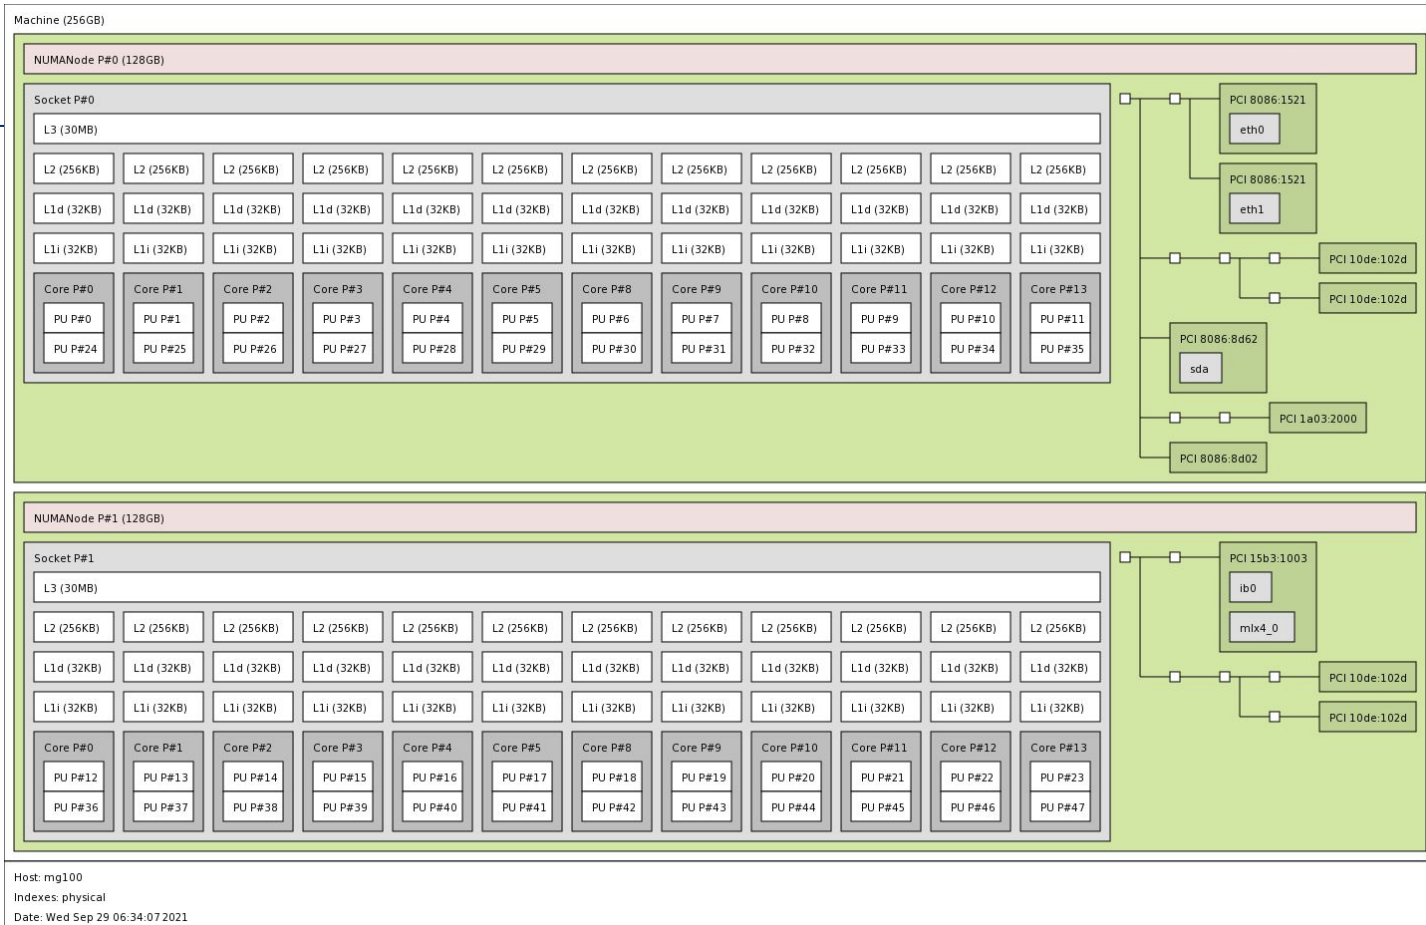
You won't need to care about the storage complexity after you have loaded your data. But data movements within a single compute node may still hold you back.

Let's have a look at a typical compute node with GPUs to accelerate training.

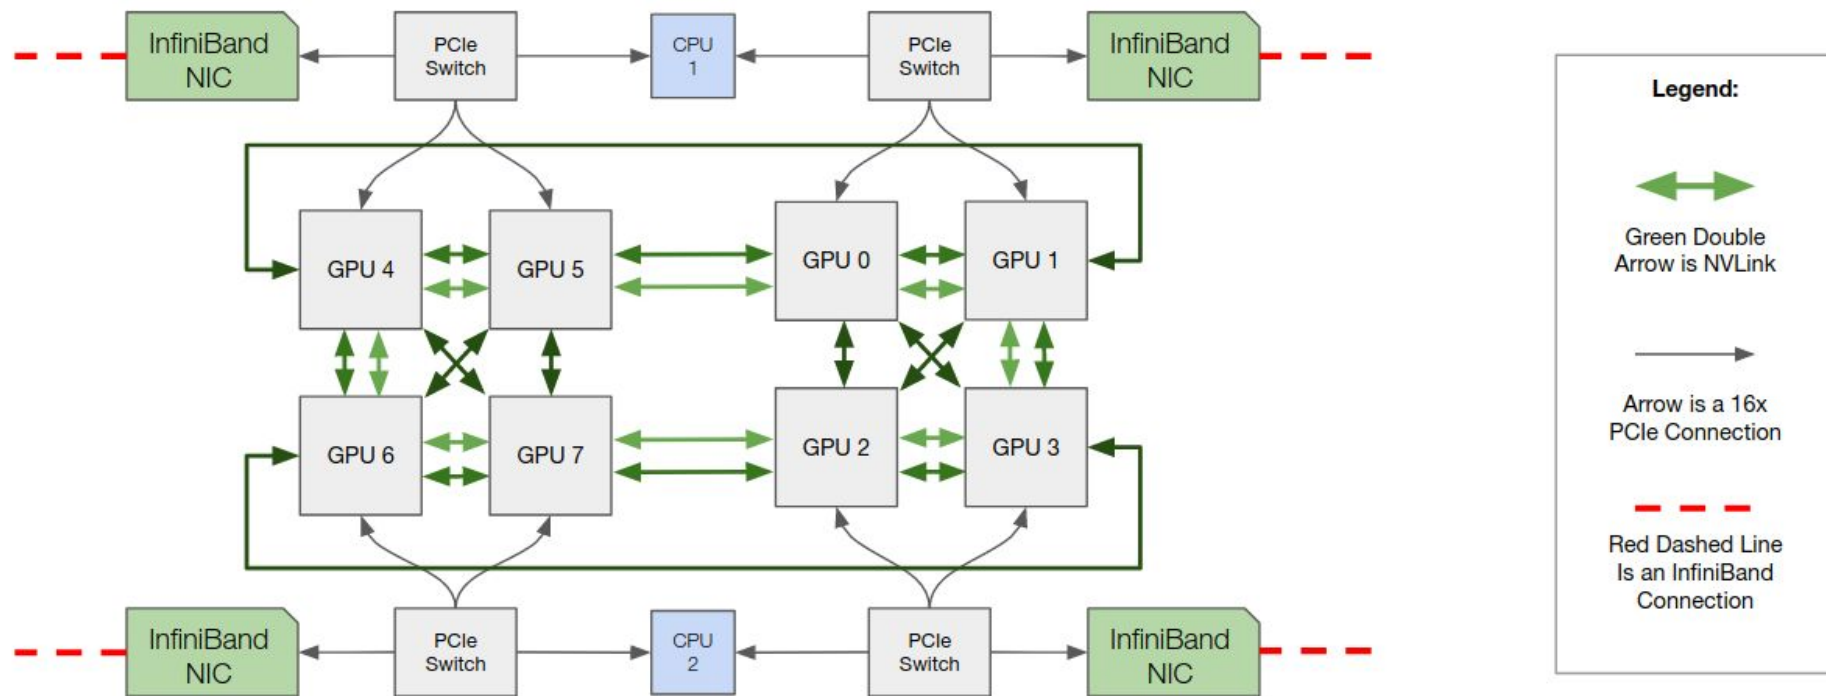


**An Example of Compute Node Architecture**

# HPC Nodes



# A top of the line GPU node: ~42 GB/s in/out of node 600 GB/s between GPUs within node

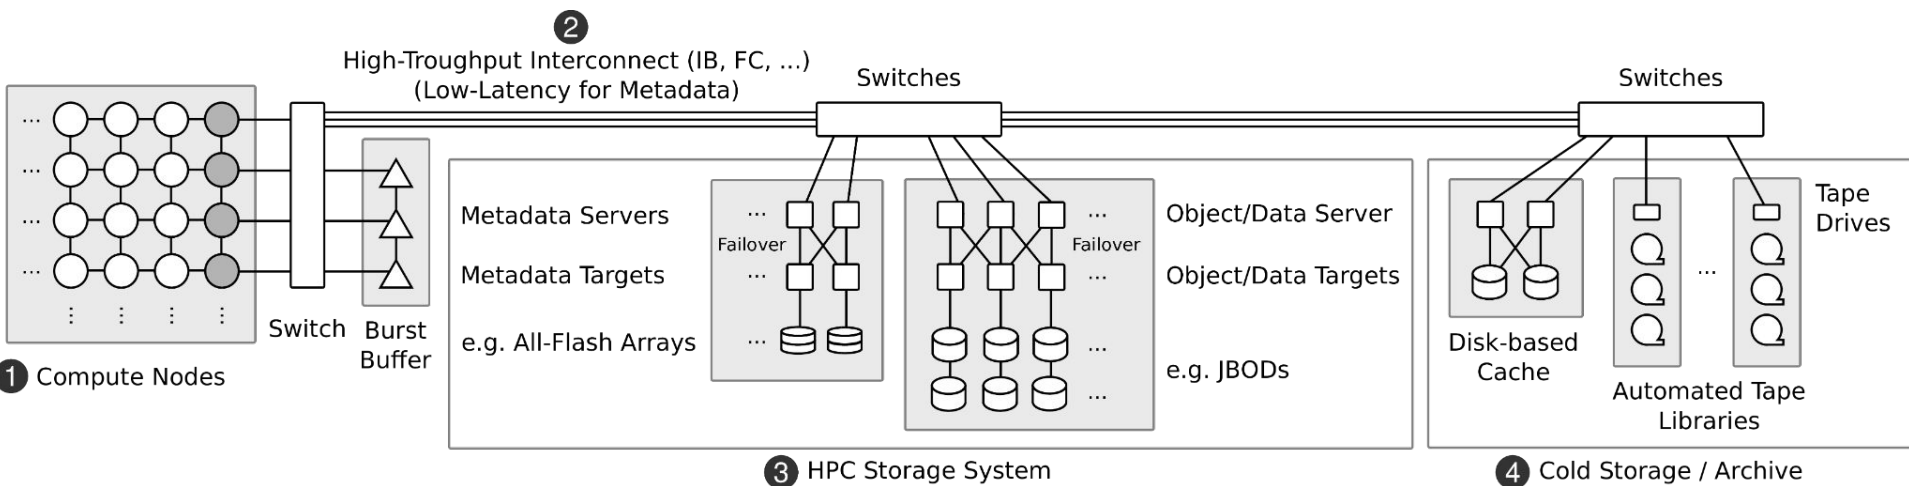


# HPC Storage?

```
[luettgau1@jwlogin06 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        378G   0    378G   0% /dev
tmpfs           378G  195M   377G   1% /dev/shm
tmpfs           378G  4.1G   374G   2% /run
tmpfs           378G   0    378G   0% /sys/fs/cgroup
/dev/mapper/centos-home 29G   4.8G   23G  18% /
/dev/md0        990M  137M   787M  15% /boot
/dev/sdd2       200M   0    200M   0% /boot/efi_disk2
/dev/sdc1       200M  6.9M   193M   4% /boot/efi
/dev/mapper/centos-tmp  20G   59M   19G   1% /tmp
/dev/mapper/centos-var  48G  2.0G   44G   5% /var
arch            1.1P  395T   630T  39% /p/arch
fastdata      8.6P  7.2P  1.4P  85% /p/fastdata
project      4.3P  2.5P  1.9P  57% /p/project
usersoftware    33T  423G   32T   2% /p/usersoftware
home         61T  14T  48T  22% /p/home
largedata_restore 4.9P  4.5P  445T  92% /p/largedata_restore
juwels_software  17T  4.4T   12T  27% /p/software/juwels
largedata2      8.0P  1.5P   6.5P  19% /p/largedata2
arch2           1021T 388T   634T  38% /p/arch2
largedata       27P   22P   4.9P  82% /p/largedata
scratch         13P   6.8P   6.0P  54% /p/scratch
```



# System: Hardware Perspective



## Node Topology as exposed by Slurm:

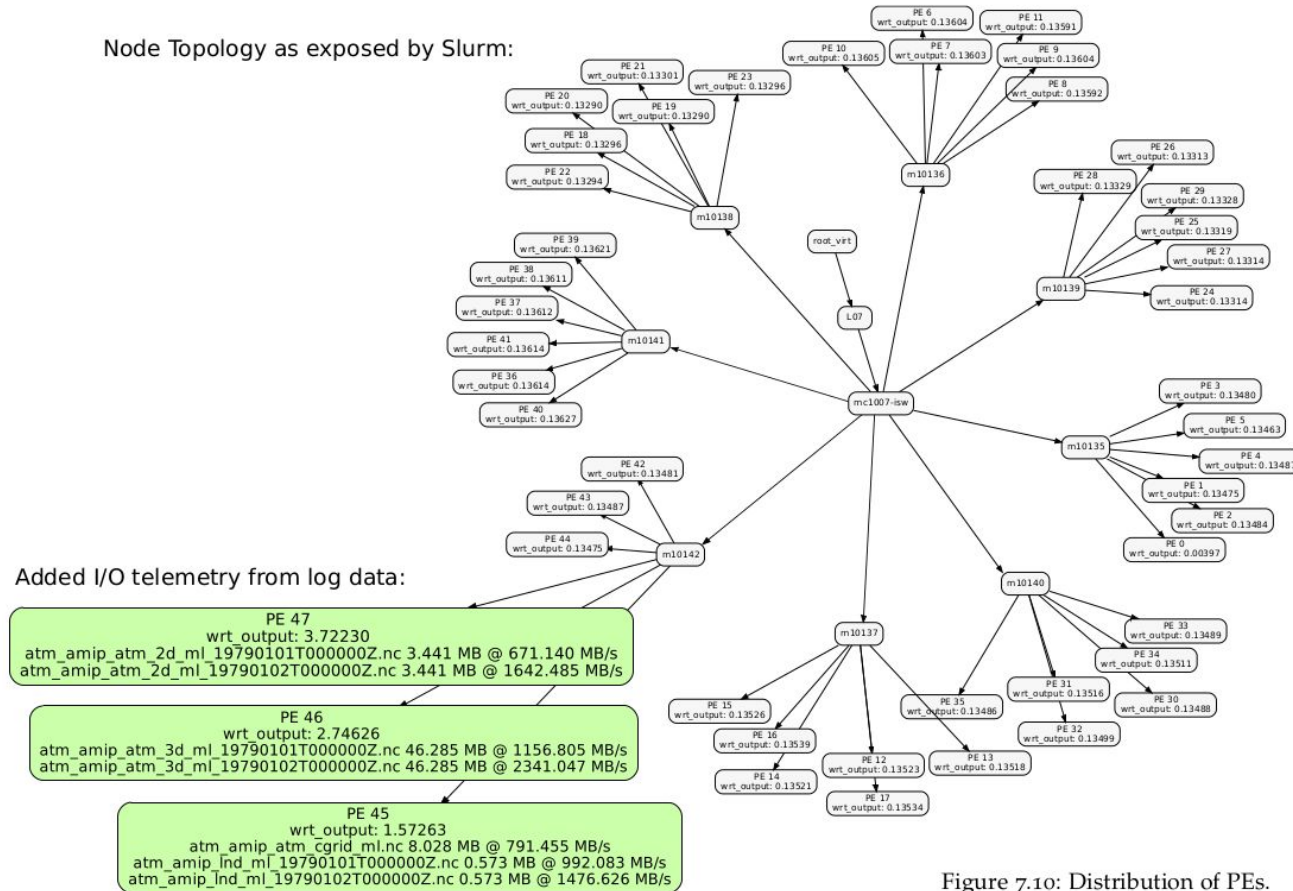
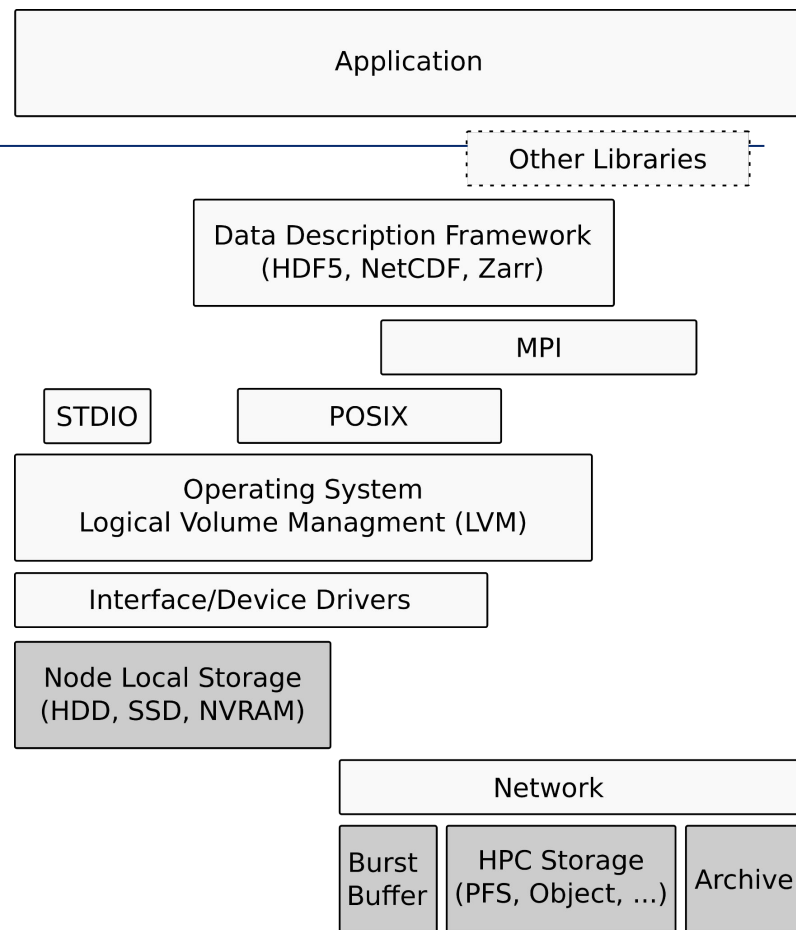


Figure 7.10: Distribution of PEs.  
Note the I/O PEs responsible for writing to the parallel file system.

# System: Software Stack

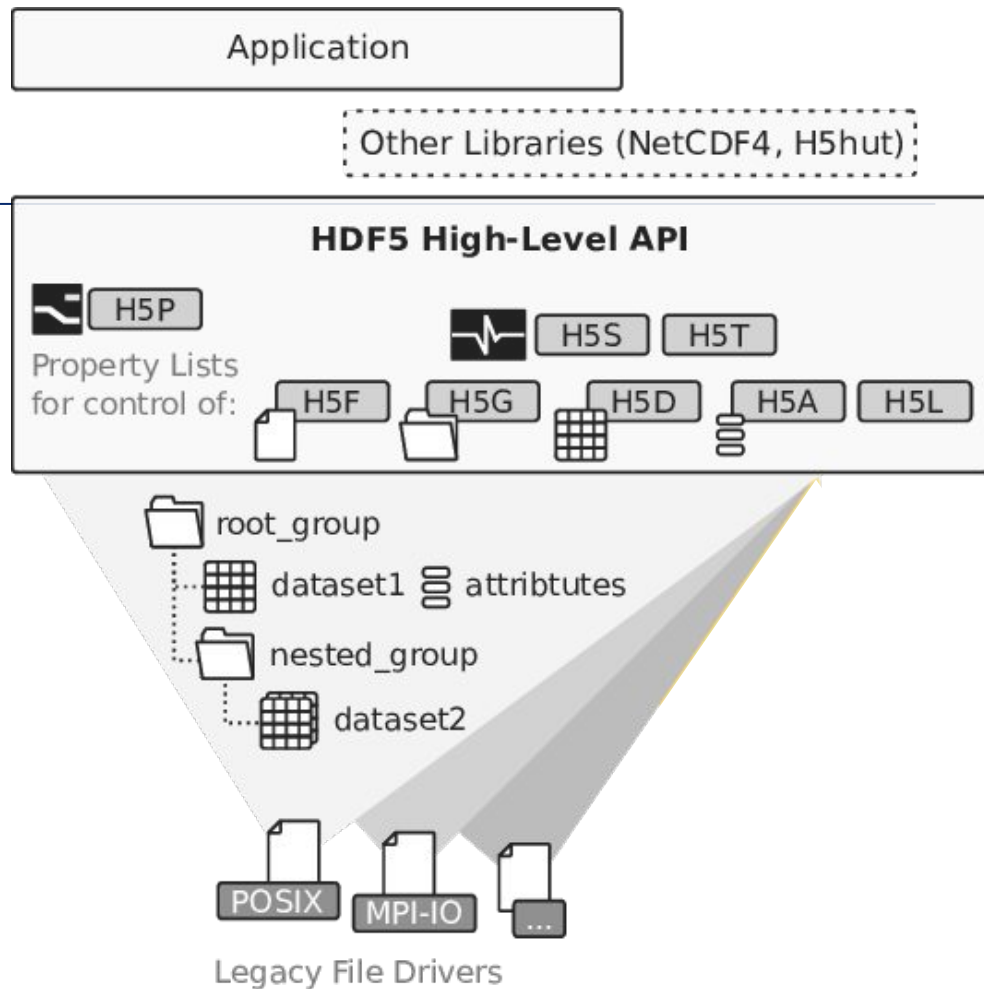


---

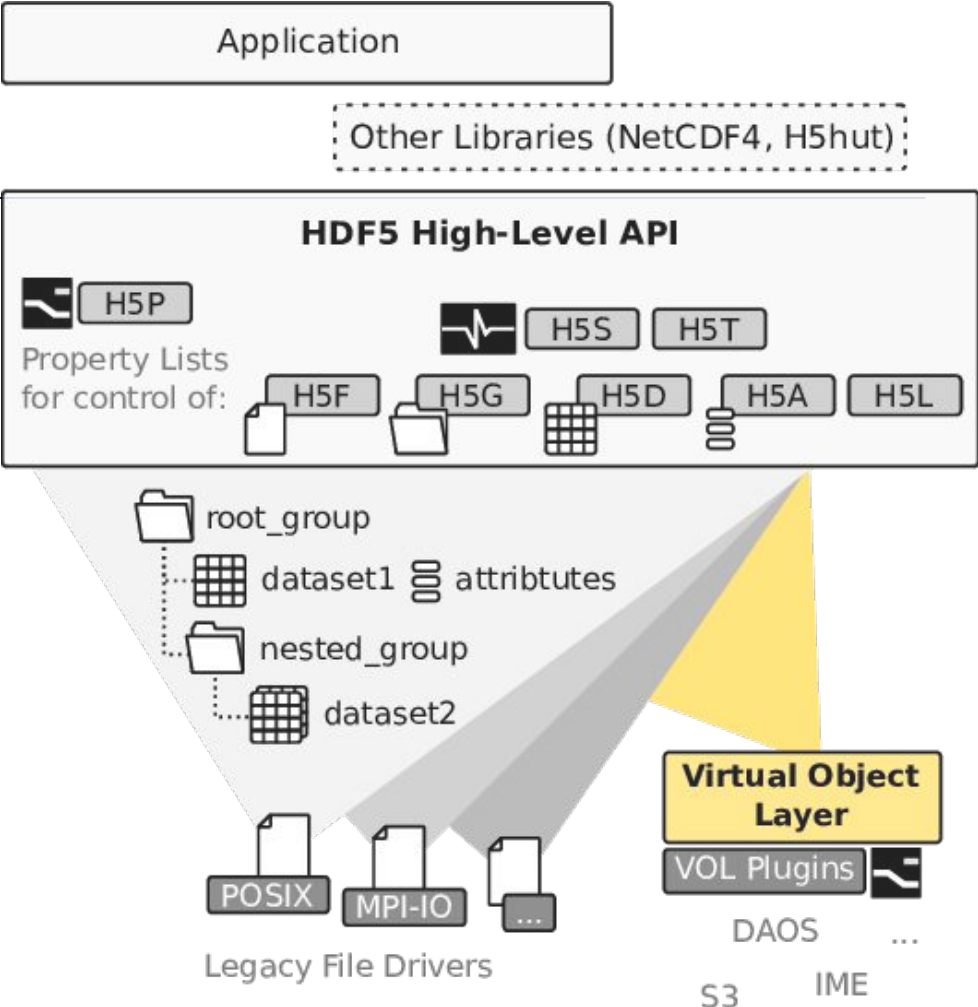
# Self-Describing Data Formats

HDF5, NetCDF, Zarr, ...

# HDF5



# HDF5



## Application

Proc Proc Proc Proc ...

dataset: {ndims, size[ndims]}

hyperslab: {count[ndims], offset[ndims],  
stride[ndims], blocks[ndims]}

chunksize: {size[ndims]}

## HDF5 High-Level API



H5P

Property Lists  
for control of:



H5S



H5T



H5F



H5G



H5D



H5A



root\_group



dataset1



attributes

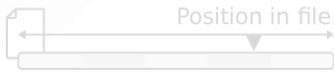


nested\_group

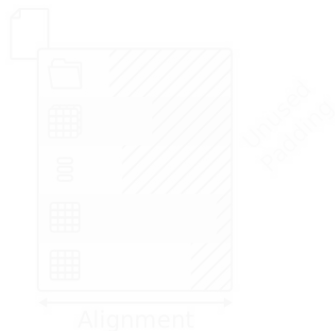


dataset2

MPI-IO  
POSIX



Alignment



© 2015-2016 Helmholtz AI

## Application

Proc Proc Proc Proc ...

dataset: {ndims, size[ndims]}

hyperslab: {count[ndims], offset[ndims],  
stride[ndims], blocks[ndims]}

chunksize: {size[ndims]}

## HDF5 High-Level API



H5P

Property Lists  
for control of:



H5S



H5T



H5F



H5G



H5D



H5A



root\_group



dataset1



attributes



nested\_group

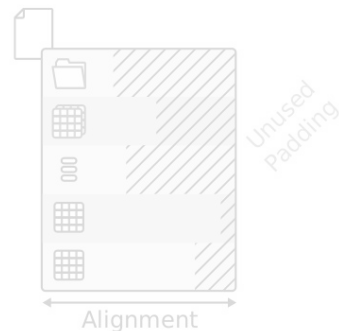


dataset2

MPI-IO  
POSIX



Alignment



Unused  
Padding

Alignment

Storage Optimization from Application, over HDF5, to file I/O targeting Lustre



## Application

Proc Proc Proc Proc ...

dataset: {ndims, size[ndims]}

hyperslab: {count[ndims], offset[ndims],  
stride[ndims], blocks[ndims]}

chunksize: {size[ndims]}

## HDF5 High-Level API



H5P

Property Lists  
for control of:



H5S



H5T



H5F



H5G



H5D



H5A



root\_group



dataset1



attributes

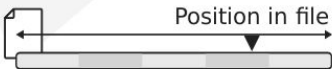


nested\_group

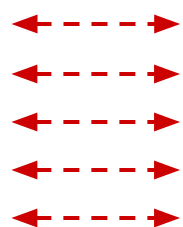
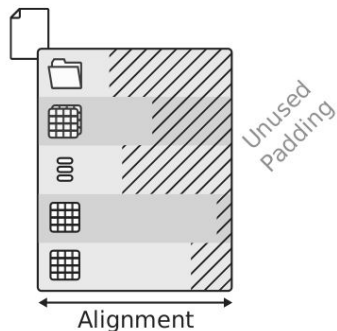


dataset2

MPI-IO  
POSIX



Alignment



## Application

Proc Proc Proc Proc ...

dataset: {ndims, size[ndims]}

hyperslab: {count[ndims], offset[ndims],  
stride[ndims], blocks[ndims]}

chunksize: {size[ndims]}

## HDF5 High-Level API



H5P

Property Lists  
for control of:



H5S



H5T



H5F



H5G



H5D



H5A

root\_group



dataset1



attributes



nested\_group

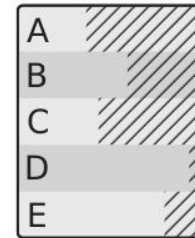
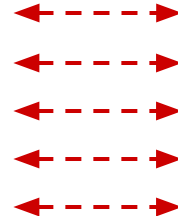
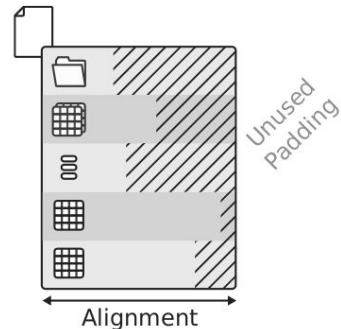


dataset2

MPI-IO  
POSIX



Alignment



Stripesize = Alignment  
Stripecount = #Targets



A D B E C

Storage Optimization from Application, over HDF5, to file I/O targeting Lustre

## Application

Proc Proc Proc Proc ...

dataset: {ndims, size[ndims]}

hyperslab: {count[ndims], offset[ndims],  
stride[ndims], blocks[ndims]}

chunksize: {size[ndims]}

## HDF5 High-Level API



H5P

Property Lists  
for control of:



H5S



H5T



H5F



H5G



H5D



H5A

root\_group



dataset1



attributes



nested\_group

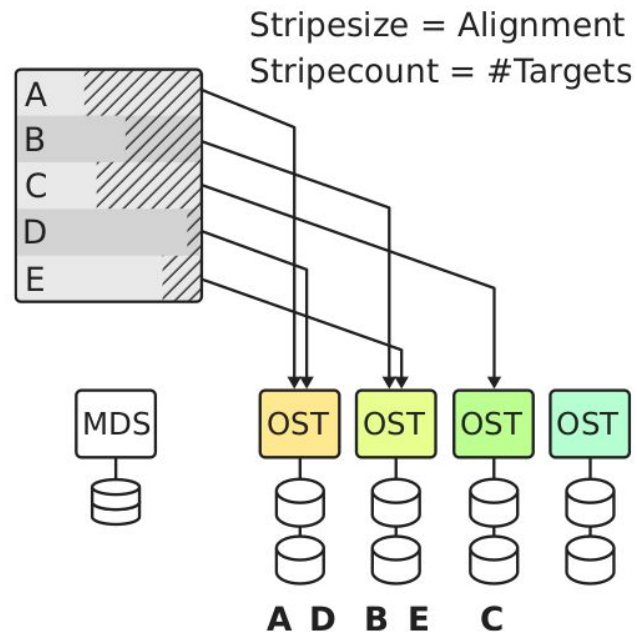
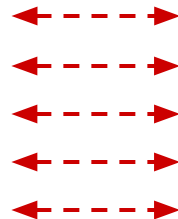
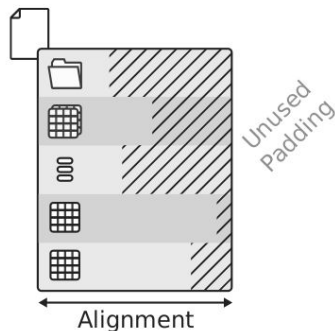


dataset2

MPI-IO  
POSIX



Alignment



Storage Optimization from Application, over HDF5, to file I/O targeting Lustre

## Application

Proc Proc Proc Proc ...

dataset: {ndims, size[ndims]}

hyperslab: {count[ndims], offset[ndims],  
stride[ndims], blocks[ndims]}

chunksize: {size[ndims]}

## HDF5 High-Level API



H5P

Property Lists  
for control of:



H5S



H5T



H5F



H5G



H5D



H5A

root\_group



dataset1



attributes

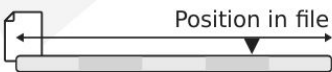


nested\_group

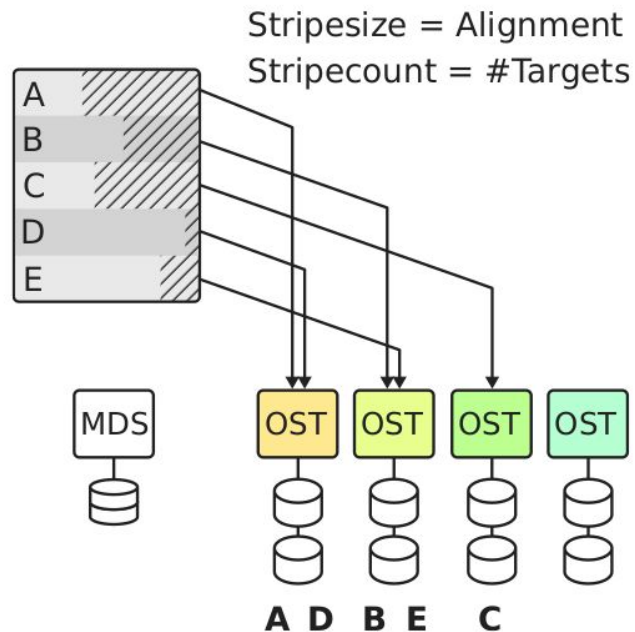
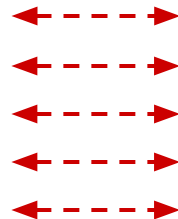
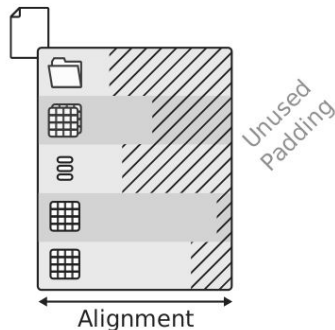


dataset2

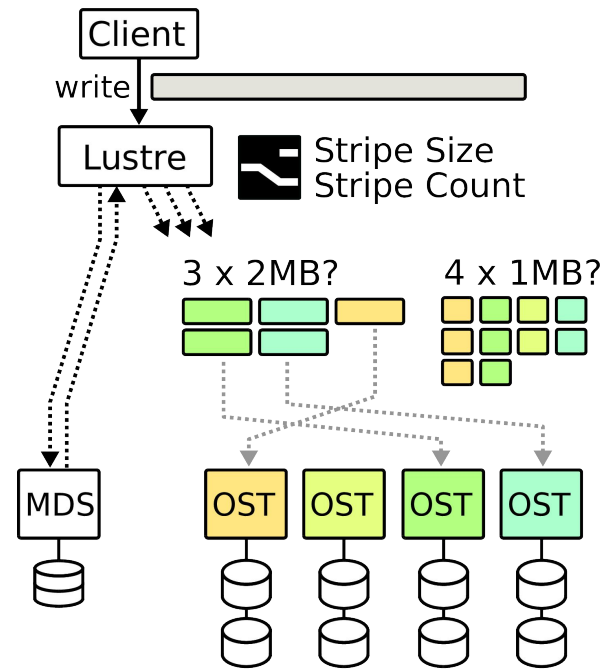
MPI-IO  
POSIX



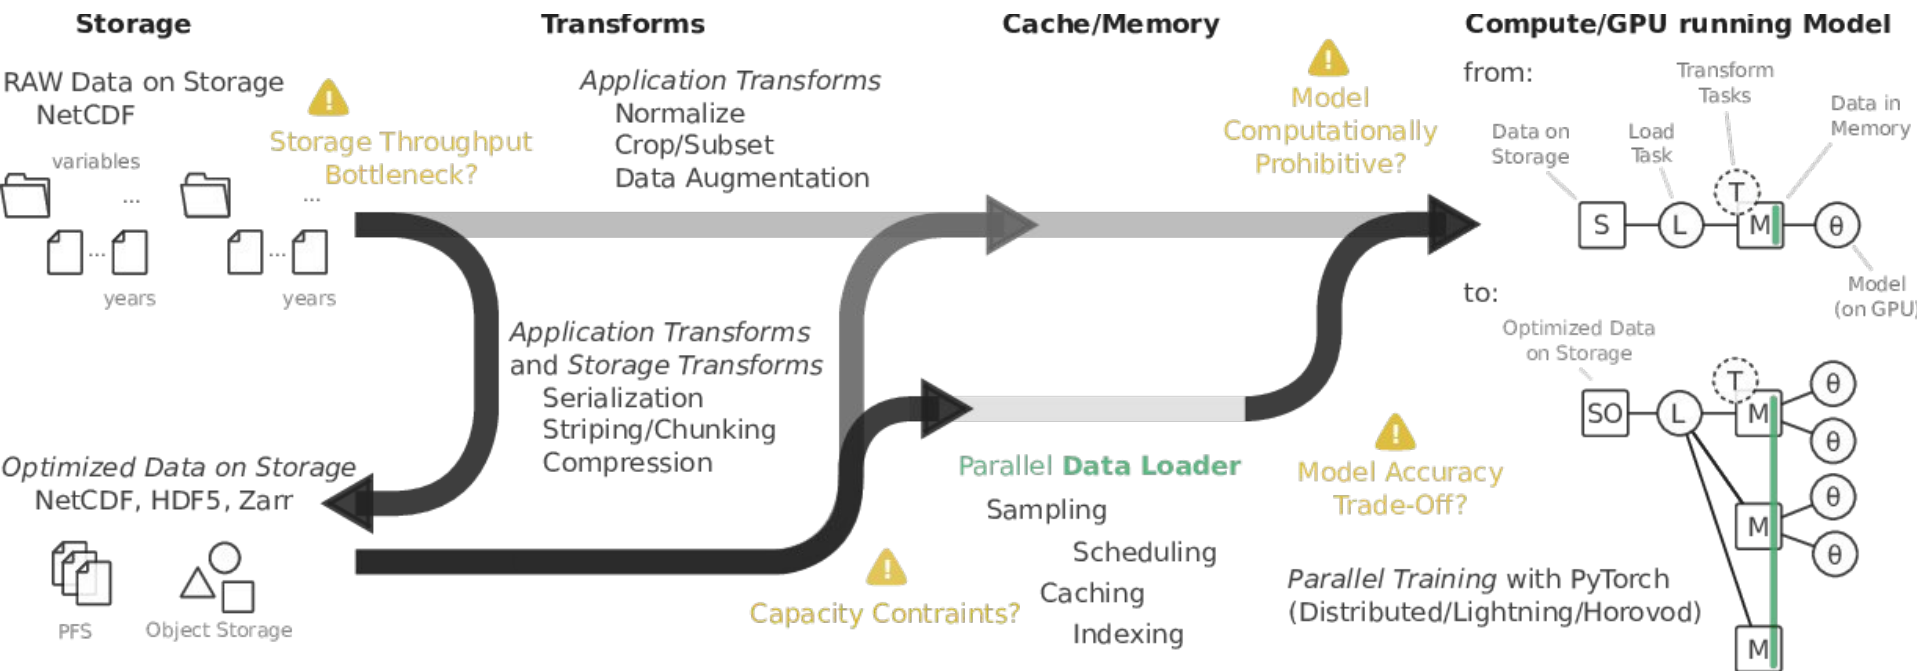
Alignment



Storage Optimization from Application, over HDF5, to file I/O targeting Lustre



# The Pre-Processing Workflow Revisited



# BREAK (10 min)

---

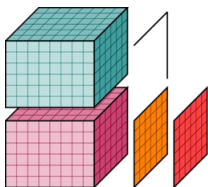


---

# Convenient Interface to Gridded Data in Python

## xarray





- Scientific data is often labeled
- Python formats
  - numpy: efficient but hard to interpret
  - Python pandas: self-explained but not performant
- xarray provides labeled arrays & datasets
- Very useful for working with netCDF
- Integrates with dask for parallel computing (which we do not cover today)

```
[4]: url = "https://opendata.arcgis.com/datasets/dd4580c810204019a7b8eb3e0b329dd6_0.csv"
      %time data = pd.read_csv(url)
      Wall time: 57 s
```

```
[13]: data
```

	ObjectId	IdBundesland	Bundesland	Landkreis	Altersgruppe	Geschlecht	AnzahlFall	AnzahlTodesfall	Meldedatum
0	1	1	Schleswig-Holstein	SK Flensburg	A00-A04	M	1	0	2020/09/30 00:00:00+00
1	2	1	Schleswig-Holstein	SK Flensburg	A00-A04	M	1	0	2020/10/29 00:00:00+00
2	3	1	Schleswig-Holstein	SK Flensburg	A00-A04	M	1	0	2020/11/03 00:00:00+00
3	4	1	Schleswig-Holstein	SK Flensburg	A00-A04	M	1	0	2020/11/20 00:00:00+00
4	5	1	Schleswig-Holstein	SK Flensburg	A00-A04	M	1	0	2020/11/23 00:00:00+00
...	...	...	...	...	...	...	...	...	...
2400782	2400783	16	Thüringen	Altenburger Land	A80+	W	1	0	2021/08/07 00:00:00+00
2400783	2400784	16	Thüringen	Altenburger Land	A80+	W	1	0	2021/08/24 00:00:00+00

```
[15]: data.values[:5]
```

```
[15]: array([[1, 1, 'Schleswig-Holstein', 'SK Flensburg', 'A00-A04', 'M', 1, 0,
        '2020/09/30 00:00:00+00', 1001, '28.09.2021, 00:00 Uhr', 0, -9,
        '2020/09/30 00:00:00+00', 0, 1, 0, 'Nicht übermittelt'],
        [2, 1, 'Schleswig-Holstein', 'SK Flensburg', 'A00-A04', 'M', 1, 0,
        '2020/10/29 00:00:00+00', 1001, '28.09.2021, 00:00 Uhr', 0, -9,
        '2020/10/29 00:00:00+00', 0, 1, 0, 'Nicht übermittelt'],
        [3, 1, 'Schleswig-Holstein', 'SK Flensburg', 'A00-A04', 'M', 1, 0,
        '2020/11/03 00:00:00+00', 1001, '28.09.2021, 00:00 Uhr', 0, -9,
        '2020/11/03 00:00:00+00', 0, 1, 0, 'Nicht übermittelt'],
        [4, 1, 'Schleswig-Holstein', 'SK Flensburg', 'A00-A04', 'M', 1, 0,
        '2020/11/20 00:00:00+00', 1001, '28.09.2021, 00:00 Uhr', 0, -9,
        '2020/11/19 00:00:00+00', 0, 1, 1, 'Nicht übermittelt'],
        [5, 1, 'Schleswig-Holstein', 'SK Flensburg', 'A00-A04', 'M', 1, 0,
        '2020/11/23 00:00:00+00', 1001, '28.09.2021, 00:00 Uhr', 0, -9,
        '2020/11/18 00:00:00+00', 0, 1, 1, 'Nicht übermittelt']],
        dtype=object)
```

# Access data with xarray

---

```
import xarray

ds = xarray.open_dataset(data_file)

ds
```

## Labeled datasets

→ Panda-like access to variables

```
ds["windspeed"]

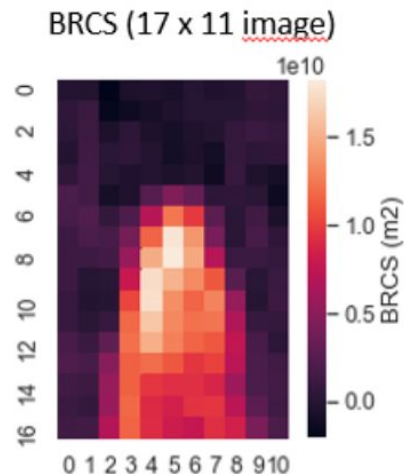
windspeed = ds.windspeed
```

# xarray in Jupyterhub

- View the dataset
- Coordinates
- Dimensions
- Data variables: type (float32, int, bool, ...), values
- 

```
[36]: ds # View the dataset
```

```
[36]: <xarray.Dataset>
Dimensions:          (delay: 17, doppler: 11, sample: 24873)
Coordinates:
  * sample            (sample) int64 0 1 2 3 4 5 ... 24868 24869 24870 24871 24872
Dimensions without coordinates: delay, doppler
Data variables:
  windspeed           (sample) float32 ...
  ddm_timestamp       (sample) float32 ...
  brcs                 (sample, delay, doppler) float32 -134927950.0 ... 756307260.0
```



# Lazy loading

---

## Lazy loading

- Data is only loaded in memory on request
- Computations etc. can be conducted without loading data
- Useful for large datasets!

## Example: unit transformation

- Define the arithmetic operation
- Only when data is loaded, it will be executed

```
windspeed = ds.windspeed  
  
# this would not load the data  
windspeed_kmh = 3.6 * windspeed  
  
# this would load the data  
windspeed_kmh.values
```

# Data cleaning with xarray

---

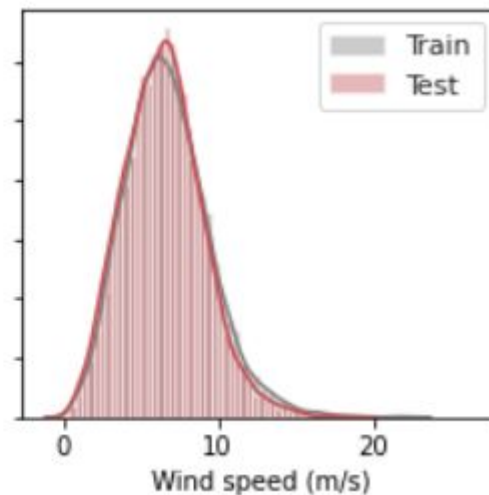
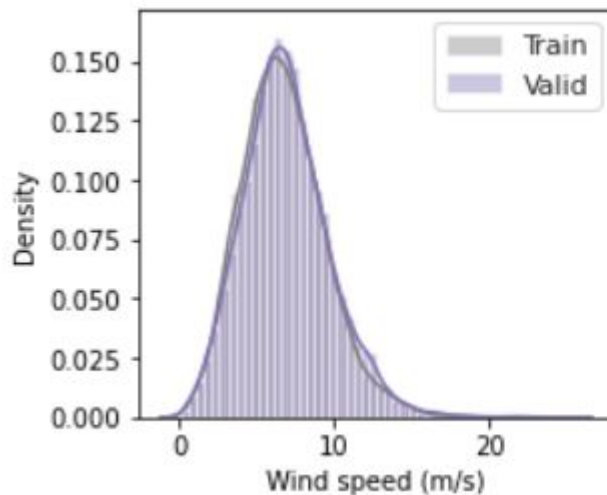
- Create a mask to select samples with None values
- Drop the samples

```
mask = xarray.ufuncs.isnan(ds_train.windspeed)
mask
```

```
ds_train = ds_train.sel(sample=~mask, drop=True)
ds_train
```

# Train / Validation / Test Dataset

- Divide before improving on the ML algorithm
- Final evaluation is done on a test set - typically in distribution!
- Check the distribution of features and labels



# PyTorch Dataset

- Load preprocessed data
- Dataset holds samples (features and labels)

```
from torch.utils.data import Dataset, DataLoader
```

- Object oriented programming
- Your dataset class is a subclass of Dataset
- In subclass, overwrite
  - `__len__`
  - `__getitem__`

```
class CyGNSSDataset(Dataset):
    def __init__(self, flag):
        '''
        Load data from hdf5 file

        Parameters:
        -----
        flag : string
            Any of train / valid / test. Defines dataset.
        -----
        Returns: dataset
        '''
        self.h5_file = h5py.File(flag + '_data.h5', 'r')

        self.y = self.h5_file['windspeed'][:].astype(np.float32)
        self.X = self.h5_file['brcs'][:].astype(np.float32)

        print(f'load {flag} input data: {self.X.shape} ({self.X.nbytes // 1e6}MB)')
        print(f'load {flag} labels: {self.y.shape} ({self.y.nbytes // 1e6}MB)')

    def __len__(self):
        '''required function for the pytorch dataloader: returns len(samples)'''
        return self.X.shape[0]

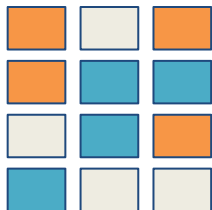
    def __getitem__(self, idx):
        '''required function for the pytorch dataloader: yields sample at idx'''
        X = self.X[idx]
        y = self.y[idx]
        return (X, y)
```

# PyTorch Dataloader

Dataset with  $N = 12$  samples

Batch size  $k = 4$

Shuffled minibatches



Dataloader retrieves the minibatch:

1st batch 

2nd batch 

3rd batch 

The DataLoader is a Python *iterable*

- Returns its members one at a time
- Implements `__iter__()`
- Implements `__getitem__()` → Dataset
- Use e.g. in for-loops

```
from torch.utils.data import DataLoader

train_loader = DataLoader(train_dataset)

for i, sample in enumerate(train_loader):
    print(i, sample)
```

0, 



---

# Measuring Performance

## Timing and Counting

# Measuring Performance

---

Consider the following example?

What are some information you might care about to plan resource usage or efficiency of your pre-processing workflow?

```
with open('filename.txt') as fp:
    for line in fp:
        data.append(line)
```

- 1) Latency: Time elapsed / wallclock time
- 2) Amount: Amount of data process
- 3) Throughput
  - a) Bytes processed per second
  - b) Operations per second

# Measuring Performance

---

Consider the following example?

What are some information you might care about to plan resource usage or efficiency of your pre-processing workflow?

```
with open('filename.txt') as fp:  
    for line in fp:  
        data.append(line)
```

# Measuring Performance

---

```
from timeit import default_timer

nbytes = 0

start = default_timer()

with open('filename.txt') as fp:
    for line in fp:
        data.append(line)

        nbytes += len(line)

end = default_timer()

elapsed_time = end - start
throughput = nbytes / elapsed_time
```

- 1) Latency: Time elapsed / wallclock time
- 2) Amount: Amount of data process
- 3) Throughput
  - a) Bytes processed per second
  - b) Operations per second

# Measuring time in IPython and Jupyter notebooks

---

- Line magic: `%time <python command>`
- Cell magic: `%%time`

```
%%time  
ds_train_hdf5['brcs'][:];
```

Execute a command several times, and calculate the average runtime: `%%timeit`